



FLIP ROBO

Micro Credit Project



Submitted by:

Surjit Singh

ACKNOWLEDGMENT

- ❖ First of all, I would like to thank all my mentors in Data Trained and FlipRobo Technologies for this opportunity.
- ❖ Most of the concepts used to predict the Micro Credit are learned from Data Trained Institute.
- ❖ Here I would be thankful that I got this chance to do the project, this gave me good knowledge about the data collection and model building ie., prediction of the data.

INTRODUCTION

- A Microfinance Institution (MFI) is an organization that offers financial services to low-income populations. MFS becomes very useful when targeting especially the unbanked poor families living in remote areas with not much sources of income. The Microfinance services (MFS) provided by MFI are Group Loans, Agricultural Loans, Individual Business Loans and so on.
- Many microfinance institutions (MFI), experts and donors are supporting the idea of using mobile financial services (MFS) which they feel are more convenient and efficient, and cost saving, than the traditional high-touch model used since long for the purpose of delivering microfinance services. Though, the MFI industry is primarily focusing on low-income families and are very useful in such areas, the implementation of MFS has been uneven with both significant challenges and successes.
- Today, microfinance is widely accepted as a poverty-reduction tool, representing \$70 billion in outstanding loans and a global outreach of 200 million clients.
- We are working with one such client that is in Telecom Industry. They are a fixed wireless telecommunications network provider. They have launched various products and have developed its business and organization based on the budget operator model, offering better products at Lower Prices to all value conscious customers through a strategy of disruptive innovation that focuses on the subscriber.
- They understand the importance of communication and how it affects a person's life, thus, focusing on providing their services and products to low-income families and poor customers that can help them in the need of hour.
- They are collaborating with an MFI to provide micro-credit on mobile balances to be paid back in 5 days. The Consumer is believed to be defaulter if he deviates from the path of paying back the loaned amount within the time duration of 5 days. For the loan amount of 5 (in Indonesian Rupiah), payback amount should be 6 (in Indonesian Rupiah), while, for the loan amount of 10 (in Indonesian Rupiah), the payback amount should be 12 (in Indonesian Rupiah)

Analytical Problem Framing

- ❖ Here our dataset has 209593 rows and 37 columns, using this dataset we will be building the model followed by training the data and then finally the model is tested by using 67% of the training data and 33% of the testing data.
- ❖ Since we have no null values from the dataset during the data collection stage, we can expect outliers and un-realistic values for certain variables.

	A	B	C	D	E	F	G	H	I	J
1	Variable	Definition	Comment							
2	label	Flag indicating whether the user paid back the credit amount within 5 days of issuing the loan{1:success, 0:failure}								
3	msisdh	mobile number of user								
4	aon	age on cellular network in days								
5	daily_decr30	Daily amount spent from main account, averaged over last 30 days (in Indonesian Rupiah)								
6	daily_decr90	Daily amount spent from main account, averaged over last 90 days (in Indonesian Rupiah)								
7	rental30	Average main account balance over last 30 days	Unsure of given definition							
8	rental90	Average main account balance over last 90 days	Unsure of given definition							
9	last_rech_date_ma	Number of days till last recharge of main account								
10	last_rech_date_da	Number of days till last recharge of data account								
11	last_rech_amt_ma	Amount of last recharge of main account (in Indonesian Rupiah)								
12	cnt_ma_rech30	Number of times main account got recharged in last 30 days								
13	fr_ma_rech30	Frequency of main account recharged in last 30 days	Unsure of given definition							
14	sumamnt_ma_rech30	Total amount of recharge in main account over last 30 days (in Indonesian Rupiah)								
15	medianamnt_ma_rech30	Median of amount of recharges done in main account over last 30 days at user level (in Indonesian Rupiah)								
16	medianmarechprebal30	Median of main account balance just before recharge in last 30 days at user level (in Indonesian Rupiah)								
17	cnt_ma_rech90	Number of times main account got recharged in last 90 days								
18	fr_ma_rech90	Frequency of main account recharged in last 90 days	Unsure of given definition							
19	sumamnt_ma_rech90	Total amount of recharge in main account over last 90 days (in Indonesian Rupiah)								
20	medianamnt_ma_rech90	Median of amount of recharges done in main account over last 90 days at user level (in Indonesian Rupiah)								
21	medianmarechprebal90	Median of main account balance just before recharge in last 90 days at user level (in Indonesian Rupiah)								
22	cnt_da_rech30	Number of times data account got recharged in last 30 days								
23	fr_da_rech30	Frequency of data account recharged in last 30 days								
24	cnt_da_rech90	Number of times data account got recharged in last 90 days								
25	fr_da_rech90	Frequency of data account recharged in last 90 days								
26	cnt_loans30	Number of loans taken by user in last 30 days								
27	amnt_loans30	Total amount of loans taken by user in last 30 days								
28	maxamnt_loans30	maximum amount of loan taken by the user in last 30 days	There are only two options: 5 & 10 Rs., for which the user needs to pay back 6 & 12 Rs. respectively							
29	medianamnt_loans30	Median of amounts of loan taken by the user in last 30 days								
30	cnt_loans90	Number of loans taken by user in last 90 days								
31	amnt_loans90	Total amount of loans taken by user in last 90 days								
32	maxamnt_loans90	maximum amount of loan taken by the user in last 90 days								
33	medianamnt_loans90	Median of amounts of loan taken by the user in last 90 days								
34	payback30	Average payback time in days over last 30 days								
35	payback90	Average payback time in days over last 90 days								
36	pcircle	telecom circle								
37	pdate	date								

Analysis

Importing the Required libraries :

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

Data Collection:

```
In [2]: data = pd.read_csv("Micro credit project.csv")
data
```

```
Out[2]:
```

	Unnamed: 0	label	msisdn	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	...	maxamnt_loans30	mec
0	1	0	21408170789	272.0	3055.050000	3065.150000	220.13	260.13	2.0	0.0	...	6.0	
1	2	1	76462170374	712.0	12122.000000	12124.750000	3691.26	3691.26	20.0	0.0	...	12.0	
2	3	1	17943170372	535.0	1398.000000	1398.000000	900.13	900.13	3.0	0.0	...	6.0	
3	4	1	55773170781	241.0	21.228000	21.228000	159.42	159.42	41.0	0.0	...	6.0	
4	5	1	03813182730	947.0	150.619333	150.619333	1098.90	1098.90	4.0	0.0	...	6.0	
...
209588	209589	1	22758185348	404.0	151.872333	151.872333	1089.19	1089.19	1.0	0.0	...	6.0	
209589	209590	1	95583184455	1075.0	36.936000	36.936000	1728.36	1728.36	4.0	0.0	...	6.0	
209590	209591	1	28556185350	1013.0	11843.111670	11904.350000	5861.83	8893.20	3.0	0.0	...	12.0	
209591	209592	1	59712182733	1732.0	12488.228330	12574.370000	411.83	984.58	2.0	38.0	...	12.0	
209592	209593	1	65061185339	1581.0	4489.362000	4534.820000	483.92	631.20	13.0	0.0	...	12.0	

209593 rows x 37 columns

```
In [3]: data.columns
```

```
Out[3]: Index(['Unnamed: 0', 'label', 'msisdn', 'aon', 'daily_decr30', 'daily_decr90',  
             'rental30', 'rental90', 'last_rech_date_ma', 'last_rech_date_da',  
             'last_rech_amt_ma', 'cnt_ma_rech30', 'fr_ma_rech30',  
             'sumamnt_ma_rech30', 'medianamnt_ma_rech30', 'medianmarechprebal30',  
             'cnt_ma_rech90', 'fr_ma_rech90', 'sumamnt_ma_rech90',  
             'medianamnt_ma_rech90', 'medianmarechprebal90', 'cnt_da_rech30',  
             'fr_da_rech30', 'cnt_da_rech90', 'fr_da_rech90', 'cnt_loans30',  
             'amnt_loans30', 'maxamnt_loans30', 'medianamnt_loans30', 'cnt_loans90',  
             'amnt_loans90', 'maxamnt_loans90', 'medianamnt_loans90', 'payback30',  
             'payback90', 'pcircle', 'pdate'],  
             dtype='object')
```

```
In [4]: data.shape
```

```
Out[4]: (209593, 37)
```

```
In [5]: data.dtypes
```

```
Out[5]: Unnamed: 0      int64  
label                int64  
msisdn              object  
aon                 float64  
daily_decr30        float64  
daily_decr90        float64  
rental30            float64  
rental90            float64  
last_rech_date_ma    float64  
last_rech_date_da    float64  
last_rech_amt_ma     int64  
cnt_ma_rech30        int64  
fr_ma_rech30         float64  
sumamnt_ma_rech30    float64  
medianamnt_ma_rech30 float64  
medianmarechprebal30 float64  
cnt_ma_rech90        int64  
fr_ma_rech90         int64  
sumamnt_ma_rech90    int64  
medianamnt_ma_rech90 float64  
medianmarechprebal90 float64  
cnt_da_rech30        float64  
fr_da_rech30         float64  
cnt_da_rech90        int64
```

- Here, as I can see that I have the features with “Float and int” datatypes except the feature “pdate” which is “object” datatype.

Statistical Analysis of the data

```
data.describe()
```

	Unnamed: 0	label	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	last_r
count	209593.000000	209593.000000	209593.000000	209593.000000	209593.000000	209593.000000	209593.000000	209593.000000	209593.000000	20
mean	104797.000000	0.875177	8112.343445	5381.402289	6082.515068	2692.581910	3483.406534	3755.84780	3712.202921	
std	60504.431823	0.330519	75696.082531	9220.623400	10918.812767	4308.586781	5770.461279	53905.89223	53374.833430	
min	1.000000	0.000000	-48.000000	-93.012667	-93.012667	-23737.140000	-24720.580000	-29.000000	-29.000000	
25%	52399.000000	1.000000	246.000000	42.440000	42.692000	280.420000	300.260000	1.000000	0.000000	
50%	104797.000000	1.000000	527.000000	1469.175667	1500.000000	1083.570000	1334.000000	3.000000	0.000000	
75%	157195.000000	1.000000	982.000000	7244.000000	7802.790000	3356.940000	4201.790000	7.000000	0.000000	
max	209593.000000	1.000000	999860.755200	265926.000000	320630.000000	198926.110000	200148.110000	998650.37770	999171.809400	5

8 rows × 34 columns

```
data.describe(include = "O")
```

	msisdn	pcircle	pdate
count	209593	209593	209593
unique	186243	1	82
top	04581185330	UPW	04-07-2016
freq	7	209593	3150

Documentation:-

Data has few columns in which the difference between mean and the standard deviation is more and, in few columns, it is less and is appropriate that few columns has mean value higher than standard deviation and also there are few columns in which standard deviation is higher than the mean value and also we can see that statistical analysis of the object datatype columns also in which the unique values of the data are mentioned and also we get more information regarding the frequent values present in the data of the columns.

Dropping the unrequired columns:

```
data = data.drop(columns = ["Unnamed: 0", "msisdn"])
```

Checking the count for our label column "label":

```
label_column_count = pd.DataFrame(data["label"].value_counts())  
label_column_count
```


	label
1	183431
0	26162

I can see that the label column has imbalanced data in it and I have to balance the data.

Plotting the heatmap for null-values :-

```
sns.heatmap(data.isnull())
```

<AxesSubplot:>



Pre-processing:

```
pcircle = pd.DataFrame(data["pcircle"].value_counts())  
pcircle
```

	pcircle
UPW	209593

```
data = data.drop(columns = ["pcircle"])
```

Pre-processing of the column "pdate" :

```
data["Pdate"] = pd.to_datetime(data.pdate,format = "%d-%m-%Y").dt.day
```

```
data["Pmonth"] = pd.to_datetime(data.pdate,format = "%d-%m-%Y").dt.month
```


```
data["Pyear"] = pd.to_datetime(data.pdate,format = "%d-%m-%Y").dt.year
```

I can see that from the column "pdate", multiple columns are extracted with the help of "pd.to_datetime".

Visualization

Label:


```
plt.figure(figsize=(5,3),dpi=80)
sns.countplot(data.label);
```



I can see that the column has an attribute(non-defaulter) with very high count than the other attribute (defaulter)

Aon :


```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.boxplot(data['aon']);
plt.subplot(2,2,2)
sns.distplot(data['aon']);
```



I can see that the column has many number of outliers present and also there are dense in nature and the distribution peak is also very narrow.

Rental90 :


```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['rental90'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['rental90']);
plt.subplot(2,2,3)
sns.distplot(data['rental90']);
```



I can see that the column has many outliers and the distribution curve has the narrow peak and also has skewness

Cnt_ma_rech30:


```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['cnt_ma_rech30'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['cnt_ma_rech30']);
plt.subplot(2,2,3)
sns.distplot(data['cnt_ma_rech30']);
```



I can see that the column has many outliers and the distribution curve has the narrow peak and also has skewness.

Amnt_loans30 :


```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['amnt_loans30'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['amnt_loans30']);
plt.subplot(2,2,3)
sns.distplot(data['amnt_loans30']);
```



I can see that the column has a large number of outliers which are dense in nature and the distribution curve has multiple peaks and also is with skewness.

Medianamnt_loans :

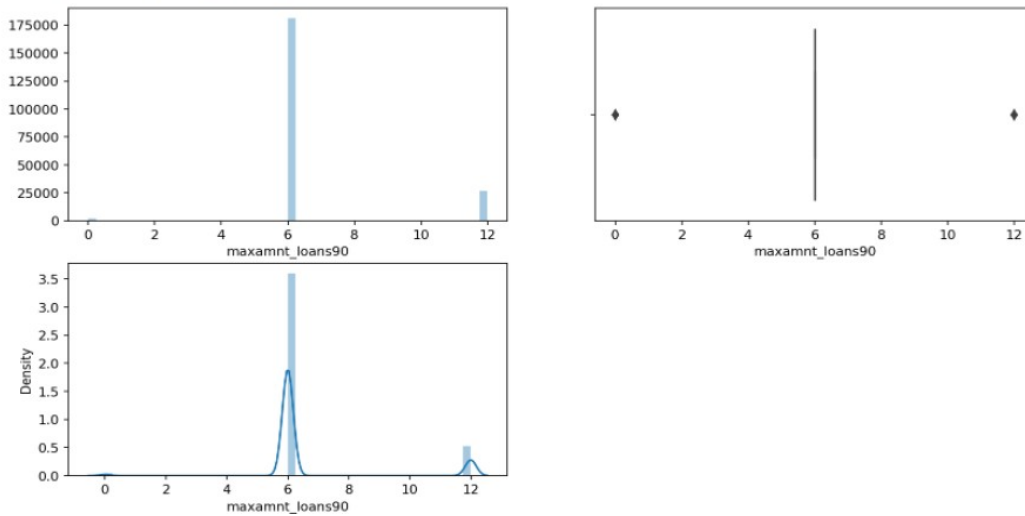
```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['medianamnt_loans30'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['medianamnt_loans30']);
plt.subplot(2,2,3)
sns.distplot(data['medianamnt_loans30']);
```



I can see that the column has very few outliers which are far away and the distribution curve is with narrow peak and also has more peaks where the data has skewness.


Maxamnt_loans90 :

```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['maxamnt_loans90'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['maxamnt_loans90']);
plt.subplot(2,2,3)
sns.distplot(data['maxamnt_loans90']);
```



I can see that the column has the outliers at the very rare end and is also very far from the quartile which can be negligible.


```
plt.figure(figsize=(5,3),dpi=150)
sns.countplot(data.maxamnt_loans90);
```



I can see that the column has the highest count for only one attribute "6" and the least count for the category "0"

Pdate :


```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['Pdate'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['Pdate']);
plt.subplot(2,2,3)
sns.distplot(data['Pdate']);
```



I can see that the column has no outliers and the distribution curve is very broad.

Pmonth :

```
plt.figure(figsize=(13,7),dpi=150)
plt.subplot(2,2,1)
sns.distplot(data['Pmonth'], kde=False);
plt.subplot(2,2,2)
sns.boxplot(data['Pmonth']);
plt.subplot(2,2,3)
sns.distplot(data['Pmonth']);
```




I can see that the column has no outliers seen and the ditribution curve is with multiple peaks.

Bivariate Analysis


➤ Aon with label

```
sns.jointplot(data=data, x='aon', y='label', kind='reg');
```



```
plt.figure(figsize = (9,5),dpi=100)  
sns.stripplot(x = 'label',y = 'aon', data = data)
```

<AxesSubplot:xlabel='label', ylabel='aon'>




I can see that the label 1 attribute has high and dense customers than the label 0 attribute.

➤ Daily_decr30 with label

```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'daily_decr30', data = data)
```

<AxesSubplot:xlabel='label', ylabel='daily_decr30'>




I can see that the high density is present in the label 1 attribute with a high value.

➤ Last_rech_date_ma with label

```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'last_rech_date_ma', data = data)
```

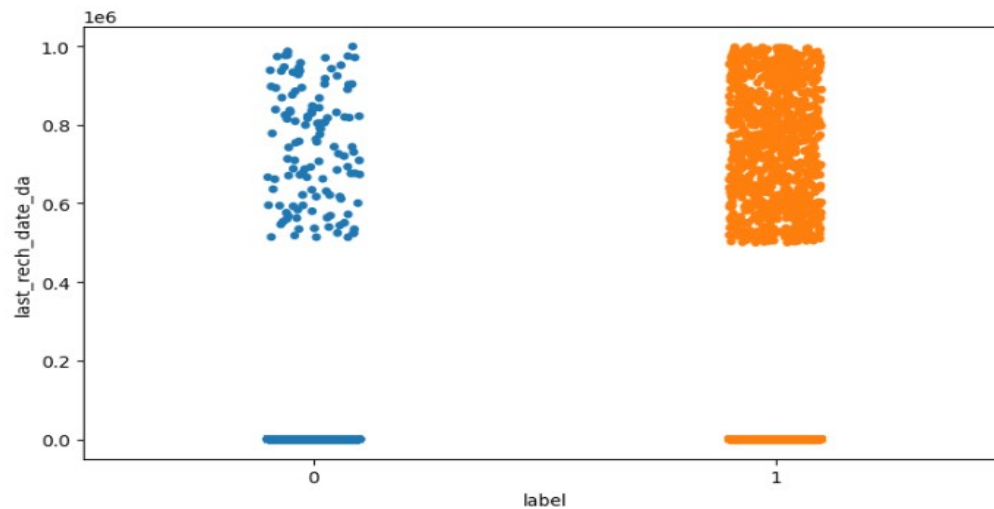
<AxesSubplot:xlabel='label', ylabel='last_rech_date_ma'>



I can see that the high density of customers who have no, of days for last recharge of main account are for label 1 attribute ie.,non-defaulters

➤ Last_rech_date_da with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'last_rech_date_da', data = data)
<AxesSubplot:xlabel='label', ylabel='last_rech_date_da'>
```



I can see that the high density of customers who have no, of days for last recharge of data account are for label 1 attribute ie.,non-defaulters.

➤ Last_rech_amt_ma with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'last_rech_amt_ma', data = data)
<AxesSubplot:xlabel='label', ylabel='last_rech_amt_ma'>
```



I can see that there is high density of customers is for both the label attributes at their starting points is same but as the amount for the last recharge increases there we can see the decrease in the density of the customers and at the highest last recharge amount point the density for both of the label attributes is almost same but when compared to label 0 ie., defaulters the label 1 ie., non-defaulters attribute has the high density.

➤ Medianamnt_ma_rech30 with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'medianamnt_ma_rech30', data = data)
<AxesSubplot:xlabel='label', ylabel='medianamnt_ma_rech30'>
```



I can see that there is high density of customers is for both the label attributes at their starting points is same but as the amount for the last recharge increases there we can see the decrease in the density of the customers and at the highest last recharge amount point the density for both of the label attributes is almost same but when compared to label 0 ie., defaulters the label 1 ie., non-defaulters attribute has the high density.

➤ Fr_ma_rech90 with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'fr_ma_rech90', data = data)
<AxesSubplot:xlabel='label', ylabel='fr_ma_rech90'>
```



I can see that the high density is present at the starting stage of the frequency of main account recharged but as there is increase in the day count the density decreased in label 0 attribute but there is density remained in the label 1 attribute and at the final point the density becomes less in label 1 attribute and it becomes negligible in label 0 attribute.

➤ Medianamnt_ma_rech90 with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'medianamnt_ma_rech90', data = data)
<AxesSubplot:xlabel='label', ylabel='medianamnt_ma_rech90'>
```



I can see that there is high density of customers is for both the label attributes at their starting points is same but as the median amount for the last recharge of the main account increases there we can see the decrease in the density of the customers and at the highest last recharge amount point the density for both of the label attributes is almost same but when compared to label 0 ie., defaulters the label 1 ie., non-defaulters attribute has the high density.

➤ Medianmarechprebal90 with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'medianmarechprebal90', data = data)
<AxesSubplot:xlabel='label', ylabel='medianmarechprebal90'>
```



I can see that the high density is present in the label 1 attribute with a high value

➤ Cnt_da_rech30 with label


```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'cnt_da_rech30', data = data)
<AxesSubplot:xlabel='label', ylabel='cnt_da_rech30'>
```



I can see that the high density of customers are for label 1 attribute ie.,non-defaulters.

➤ Payback30 with label

```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'payback30', data = data)
<AxesSubplot:xlabel='label', ylabel='payback30'>
```




I can see that the high density is present at the starting stage decreased in label 0 attribute but there is density remained in the label 1 attribute and at the final point the density becomes less in label 1 attribute and it becomes negligible in label 0 attribute.

➤ Payback90 with label

```
plt.figure(figsize = (9,5),dpi=100)
sns.stripplot(x = 'label',y = 'payback90', data = data)
```

<AxesSubplot:xlabel='label', ylabel='payback90'>



I can see that the high density is present at the starting stage decreased in label 0 attribute but there is density remained in the label 1 attribute and at the final point the density becomes less in label 1 attribute and it becomes negligible in label 0 attribute.

Correlation :

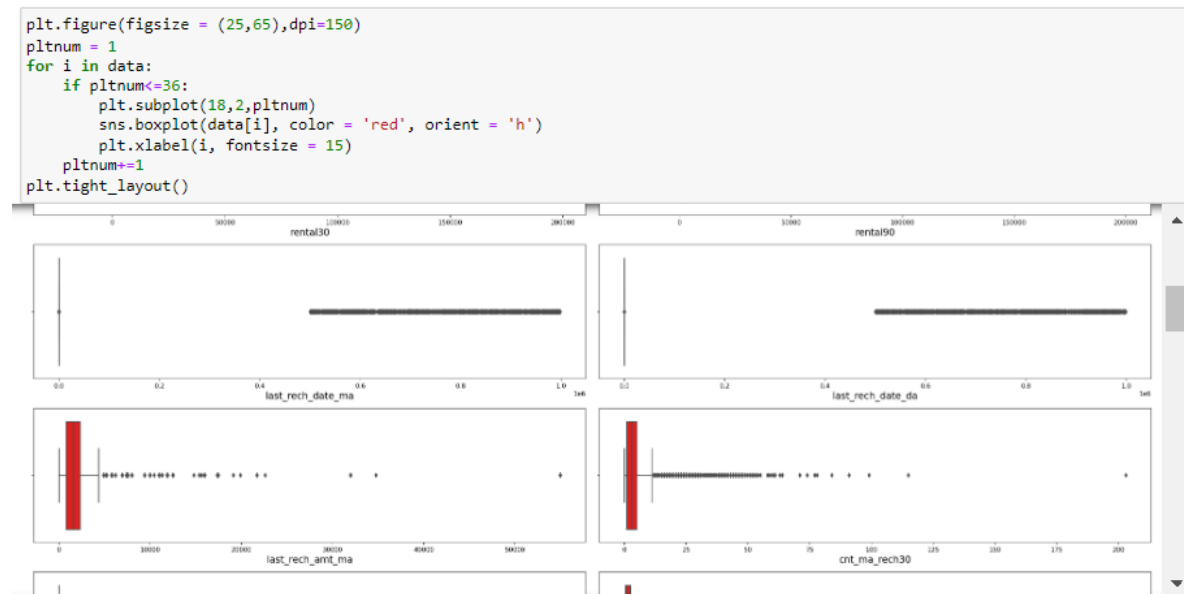
Here I find the correlation between the features and the label:

```
corr_data = data.corr()
corr_data['label'].sort_values(ascending = False)
```

label	1.000000
cnt_ma_rech30	0.237331
cnt_ma_rech90	0.236392
sumamnt_ma_rech90	0.205793
sumamnt_ma_rech30	0.202828
amnt_loans90	0.199788
amnt_loans30	0.197272
cnt_loans30	0.196283
daily_decr30	0.168298
daily_decr90	0.166150
Pmonth	0.154949
medianamnt_ma_rech30	0.141490
last_rech_amt_ma	0.131804
medianamnt_ma_rech90	0.120855
fr_ma_rech90	0.084385
maxamnt_loans90	0.084144
rental90	0.075521
rental30	0.058085
payback90	0.049183
payback30	0.048336
medianamnt_loans30	0.044589
medianmarechpreba190	0.039300
medianamnt_loans90	0.035747
Pdate	0.006825
cnt_loans90	0.004733
cnt_da_rech30	0.003827
last_rech_date_ma	0.003728
cnt_da_rech90	0.002999
last_rech_date_da	0.001711
fr_ma_rech30	0.001330
maxamnt_loans30	0.000248
fr_da_rech30	-0.000027
aon	-0.003785
medianmarechpreba130	-0.004829
fr_da_rech90	-0.005418

Name: label, dtype: float64

Detection of the Outliers



I can see that most of the columns have outliers present in them and also with dense and also with number of outliers present and so we have to treat them for better accuracy in our model building.

Treating the outliers

```
from scipy.stats import zscore
```

```
z = np.abs(zscore(data))
z.shape
```

```
(209593, 35)
```

```
threshold = 5.5
print(np.where(z>5.5))
```


```
(array([ 30, 53, 65, ..., 209531, 209533, 209576], dtype=int64), array([6, 6, 1, ..., 7, 6, 1], dtype=int64))
```

```
data_new = data[(z<5.5).all(axis = 1)]
print(data.shape)
print(data_new.shape)
```

```
(209593, 35)
(192459, 35)
```


Checking for Skewness

```
plt.figure(figsize = (50,165),dpi=150)
pltnum = 1
for i in data_new:
    if pltnum<=36:
        plt.subplot(18,2,pltnum)
        sns.distplot(data_new[i], color = 'blue')
        plt.xlabel(i, fontsize = 35)
        pltnum+=1
plt.tight_layout()
```



```
data_new.skew().sort_values()
```

label	-2.248861
Pdate	0.206862
Pmonth	0.371847
aon	0.947905
cnt_ma_rech30	1.730082
maxamnt_loans90	1.747901
cnt_ma_rech90	1.891819
cnt_loans30	1.945914
amnt_loans30	1.967762
fr_ma_rech30	2.010125
sumamnt_ma_rech30	2.176749
last_rech_amt_ma	2.221092
fr_ma_rech90	2.228096
amnt_loans90	2.244584
sumamnt_ma_rech90	2.268428
daily_decr30	2.372679
medianamnt_ma_rech30	2.451058
medianamnt_ma_rech90	2.465278
daily_decr90	2.514251
rental30	2.561126
rental90	2.689101
last_rech_date_ma	3.097659
payback90	3.601847
payback30	3.903939
medianamnt_loans30	4.075996
medianamnt_loans90	4.453088
medianmarechprebal90	5.252733
cnt_da_rech90	7.427612
last_rech_date_da	9.974328
medianmarechprebal30	10.838790
cnt_da_rech30	35.421656
maxamnt_loans30	37.890780
cnt_loans90	54.802217
fr_da_rech90	68.727574
fr_da_rech30	88.162720

dtype: float64

Most of the features have skewness and except the label column all the other feature columns are positively skewed, in that few columns are with medium positive skewness and there are also few columns with very high positive skewness.

Before removing the skewness we will split the data into features and label

```
x = data_new.drop(columns = 'label')
y = data_new['label']
```

x

	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	last_rech_amt_ma	cnt_ma_rech30	fr_ma_rech30	...	max
0	272.0	3055.050000	3065.150000	220.13	260.13	2.0	0.0	1539	2	21.0	...	
1	712.0	12122.000000	12124.750000	3691.26	3691.26	20.0	0.0	5787	1	0.0	...	
2	535.0	1398.000000	1398.000000	900.13	900.13	3.0	0.0	1539	1	0.0	...	
3	241.0	21.228000	21.228000	159.42	159.42	41.0	0.0	947	0	0.0	...	
4	947.0	150.619333	150.619333	1098.90	1098.90	4.0	0.0	2309	7	2.0	...	
...
209588	404.0	151.872333	151.872333	1089.19	1089.19	1.0	0.0	4048	3	2.0	...	
209589	1075.0	36.936000	36.936000	1728.36	1728.36	4.0	0.0	773	4	1.0	...	
209590	1013.0	11843.111670	11904.350000	5861.83	8893.20	3.0	0.0	1539	5	8.0	...	
209591	1732.0	12488.228330	12574.370000	411.83	984.58	2.0	38.0	773	5	4.0	...	
209592	1581.0	4489.362000	4534.820000	483.92	631.20	13.0	0.0	7526	2	1.0	...	

192459 rows × 34 columns

y

```
0    0
1    1
2    1
3    1
4    1
```

Scaling the data

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
scaled = scaler.fit_transform(x)
x = pd.DataFrame(scaled, columns = x.columns)
x
```

	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	last_rech_amt_ma	cnt_ma_rech30	fr_ma_rech30	...	m
0	0.128617	0.055251	0.047018	0.234328	0.187155	0.218310	0.201389	0.12312	0.074074	0.552632	...	
1	0.305466	0.216785	0.183925	0.335913	0.266956	0.345070	0.201389	0.46296	0.037037	0.000000	...	
2	0.234325	0.025730	0.021825	0.254229	0.202040	0.225352	0.201389	0.12312	0.037037	0.000000	...	
3	0.116158	0.001202	0.001019	0.232551	0.184812	0.492958	0.201389	0.07576	0.000000	0.000000	...	
4	0.399920	0.003507	0.002975	0.260046	0.206663	0.232394	0.201389	0.18472	0.259259	0.052632	...	
...
192454	0.181672	0.003529	0.002993	0.259762	0.206437	0.211268	0.201389	0.32384	0.111111	0.052632	...	
192455	0.451367	0.001481	0.001257	0.278468	0.221303	0.232394	0.201389	0.06184	0.148148	0.026316	...	
192456	0.426447	0.211817	0.180595	0.399437	0.387943	0.225352	0.201389	0.12312	0.185185	0.210526	...	
192457	0.715434	0.223310	0.190720	0.239938	0.204004	0.218310	0.465278	0.06184	0.185185	0.105263	...	
192458	0.654743	0.080805	0.069228	0.242048	0.195785	0.295775	0.201389	0.60208	0.074074	0.026316	...	

192459 rows × 34 columns

I transform the features present in the variable x to be scaled and then we will get a scaled and transformed data of x, which is used for removing the skewness present in the features.

Removing skewness

```
from sklearn.preprocessing import power_transform
```

```
transform_data = power_transform(x, method = 'yeo-johnson')  
x = pd.DataFrame(transform_data, columns = x.columns)  
x
```

	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	last_rech_amt_ma	cnt_ma_rech30	fr_ma_rech30	...
0	-0.772547	0.159357	0.095064	-0.882238	-0.891112	-0.400277	-0.124951	0.111881	-0.374333	1.857935	...
1	0.390238	1.451604	1.382771	0.947386	0.707262	1.608371	-0.124951	1.734547	-0.885060	-1.012938	...
2	-0.008536	-0.366210	-0.400540	-0.363228	-0.482142	-0.256674	-0.124951	0.111881	-0.885060	-1.012938	...
3	-0.879643	-0.935841	-0.921750	-0.933980	-0.961697	2.937812	-0.124951	-0.456250	-1.510166	-1.012938	...
4	0.815744	-0.876137	-0.867678	-0.230492	-0.367920	-0.117609	-0.124951	0.651870	1.131382	-0.090680	...
...
192454	-0.359411	-0.875566	-0.867160	-0.236795	-0.373369	-0.548592	-0.124951	1.373243	0.045864	-0.090680	...
192455	1.008365	-0.928518	-0.915124	0.141702	-0.041273	-0.117609	-0.124951	-0.656379	0.393839	-0.503695	...
192456	0.918080	1.433202	1.367321	1.446951	1.611074	-0.256674	-0.124951	0.111881	0.683766	1.223114	...
192457	1.706643	1.474640	1.412874	-0.724979	-0.432920	-0.400277	6.198337	-0.656379	0.683766	0.521957	...
192458	1.579608	0.508042	0.436928	-0.668165	-0.646009	0.955357	-0.124951	1.927979	-0.374333	-0.503695	...

192459 rows × 34 columns

Documentation :


- I can see that I have imported the “powertransform” library and passed the features present in the variable x into the powertransform and then I have removed the skewness of the data and converted into dataframe.
- Now, here we have a look at the skewness of the features whether they are changed or not.

```
x.skew()
aon 0.070230
daily_decr30 0.340121
daily_decr90 0.374143
rental30 0.363987
rental90 0.347998
last_rech_date_ma 0.647721
last_rech_date_da 6.373282
last_rech_amt_ma -0.187723
cnt_ma_rech30 -0.026788
fr_ma_rech30 0.274985
sumamnt_ma_rech30 0.108371
medianamnt_ma_rech30 -0.083915
medianmarechprebal30 0.531069
cnt_ma_rech90 0.122436
fr_ma_rech90 0.384721
sumamnt_ma_rech90 0.135732
medianamnt_ma_rech90 -0.100244
medianmarechprebal90 0.510589
cnt_da_rech30 10.540065
fr_da_rech30 74.134435
cnt_da_rech90 6.645226
fr_da_rech90 57.088248
cnt_loans30 0.306502
amnt_loans30 0.309343
maxamnt_loans30 2.234694
medianamnt_loans30 3.513169
cnt_loans90 0.246665
amnt_loans90 0.168387
maxamnt_loans90 2.154257
medianamnt_loans90 3.850225
payback30 0.224000
payback90 0.152812
Pdate -0.023090
Pmonth -0.321926
dtype: float64
```

I can see that most of the features have change in their skewness but still I have skewness present and so we use “cuberoot” from NumPy and then we pass “x” into it and assign to the variable “x” again where we can see that most of the columns have a lot of changes in their skewness except few and so we can proceed with our model building.

Balancing the data

```
sns.countplot(y)
<AxesSubplot:xlabel='label', ylabel='count'>
```




The data is imbalanced as we can see and so we will import “SMOTE” and handle the imbalanced data.

```
from imblearn.over_sampling import SMOTE
```

```
sm = SMOTE()  
x_over, y_over = sm.fit_resample(x,y)
```

```
sns.countplot(y_over)
```

```
<AxesSubplot:xlabel='label', ylabel='count'>
```



I can see that I have balanced the imbalanced data and also I can see here that both of the attributes present in the label column which were actually imbalanced are now in equal proportions and so by this I can say that our model is set for training and testing of the data.

Checking the random state

```
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score  
from sklearn.model_selection import train_test_split
```

```
rs = 0  
for i in range(0,200):  
    x_train,x_test, y_train,y_test = train_test_split(x_over,y_over,test_size = 0.33, random_state = i)  
    lg = LogisticRegression()  
    lg.fit(x_train,y_train)  
    ts_pred = lg.predict(x_test)  
    tr_pred = lg.predict(x_train)  
    ts_score = accuracy_score(y_test,ts_pred)  
    tr_score = accuracy_score(y_train, tr_pred)  
    if round(ts_score*100,1) == round(tr_score*100,1):  
        if i>rs:  
            rs = i  
print('the best random state for the data set is', rs)
```

```
the best random state for the data set is 196
```

I have used train_test_split and passed x_over, y_over which are the variables after balancing the data and we used “. fit” method to train the data and predicted the test data and accuracy score for which we got the random state as 196.

```
x_train,x_test, y_train,y_test = train_test_split(x_over,y_over,test_size = 0.33, random_state = rs)
```

I should proceed with the model testing with the testsize 33% and we present classification report and accuracy score for accuracy score.

Logistic Regression

```
logreg = LogisticRegression()  
logreg.fit(x_train,y_train)  
logreg_pred = logreg.predict(x_test)  
logreg_score = accuracy_score(y_test,logreg_pred)  
logreg_score
```

0.7769216905469947

```
print(classification_report(y_test, logreg_pred))
```

	precision	recall	f1-score	support
0	0.76	0.81	0.78	55615
1	0.80	0.74	0.77	55355
accuracy			0.78	110970
macro avg	0.78	0.78	0.78	110970
weighted avg	0.78	0.78	0.78	110970

```
print(roc_auc_score(y_test, logreg_pred))
```

0.7768460401150915

I can see that the model tested with 77% accuracy and the roc_auc_score is 77%.

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
randf = RandomForestClassifier()  
randf.fit(x_train,y_train)  
randf_pred = randf.predict(x_test)  
randf_score = accuracy_score(y_test,randf_pred)  
randf_score
```

0.9494638190501937

```
print(classification_report(y_test, randf_pred))
```

	precision	recall	f1-score	support
0	0.95	0.95	0.95	55615
1	0.95	0.95	0.95	55355
accuracy			0.95	110970
macro avg	0.95	0.95	0.95	110970
weighted avg	0.95	0.95	0.95	110970

```
print(roc_auc_score(y_test, randf_pred))
```

0.9494572919702862

I can see that the model tested with 94% accuracy and the roc_auc_score is 94%.

Extra Trees Classifier

```
from sklearn.ensemble import ExtraTreesClassifier
```

```
extr = ExtraTreesClassifier()  
extr.fit(x_train,y_train)  
extr_pred = extr.predict(x_test)  
extr_score = accuracy_score(y_test,extr_pred)  
extr_score
```

```
0.957853473911868
```

```
print(classification_report(y_test, extr_pred))
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	55615
1	0.97	0.94	0.96	55355
accuracy			0.96	110970
macro avg	0.96	0.96	0.96	110970
weighted avg	0.96	0.96	0.96	110970

```
print(roc_auc_score(y_test, extr_pred))
```

```
0.9578216775487678
```

I can see that the model tested with 95.7% accuracy and the roc_auc_score is 95.7%.

KNN Classifier

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier()  
knn.fit(x_train,y_train)  
knn_pred = knn.predict(x_test)  
knn_score = accuracy_score(y_test, knn_pred)  
knn_score
```

```
0.8740109939623322
```

```
print(classification_report(y_test, knn_pred))
```

	precision	recall	f1-score	support
0	0.81	0.97	0.89	55615
1	0.97	0.77	0.86	55355
accuracy			0.87	110970
macro avg	0.89	0.87	0.87	110970
weighted avg	0.89	0.87	0.87	110970

```
print(roc_auc_score(y_test, knn_pred))
```

```
0.8737746944433985
```

I can see that the model tested with 87% accuracy and the roc_auc_score is 87%

Checking for cross validation score

```
from sklearn.model_selection import cross_val_score
```

```
cv1 = cross_val_score(logreg, x_over,y_over,cv = 5)  
cv1 = cv1.mean()  
cv1
```

0.7766242887780394

```
cv2 = cross_val_score(randf, x_over,y_over,cv = 5)  
cv2 = cv2.mean()  
cv2
```

0.9482355209212148

```
cv3 = cross_val_score(extr, x_over,y_over,cv = 5)  
cv3 = cv3.mean()  
cv3
```

0.9641034996143866

```
cv4 = cross_val_score(knn, x_over,y_over,cv = 5)  
cv4 = cv4.mean()  
cv4
```

0.8818129466818341

I can see that out of all the models used for prediction, Extra Trees Classifier model is with high accuracy score and also high cross validation score which is 96.4%.

Model Selection

```
model =[logreg_score, randf_score, extr_score,knn_score]  
cross_val = [cv1,cv2,cv3,cv4]  
selection = pd.DataFrame({})  
selection['model'] = model  
selection['cross_val'] = cross_val  
selection['difference'] = selection['model'] - selection['cross_val']  
selection
```

	model	cross_val	difference
0	0.776922	0.776624	0.000297
1	0.949464	0.948236	0.001228
2	0.957853	0.964103	-0.006250
3	0.874011	0.881813	-0.007802

Here we can see that our best model is "Extra Trees Classifier model" with an accuracy of 96% which is highest than the rest of the models and so we choose this model for "hyper parameter tuning"

Hyper parameter tuning

```
from sklearn.model_selection import GridSearchCV
```

```
params = {'n_estimators':[0,50],  
          'criterion':['gini','entropy'],  
          'max_depth':[2,4,6],  
          'min_samples_split':[2,3,4],  
          'bootstrap':[True,False]}
```

```
final = GridSearchCV(ExtraTreesClassifier(),params,cv=5, n_jobs=-1)  
final.fit(x_train,y_train)
```

```
GridSearchCV(cv=5, estimator=ExtraTreesClassifier(), n_jobs=-1,  
             param_grid={'bootstrap': [True, False],  
                          'criterion': ['gini', 'entropy'],  
                          'max_depth': [2, 4, 6], 'min_samples_split': [2, 3, 4],  
                          'n_estimators': [0, 50]})
```

```
final.best_params_
```

```
{'bootstrap': True,  
 'criterion': 'entropy',  
 'max_depth': 6,  
 'min_samples_split': 2,  
 'n_estimators': 50}
```

```
final_rf = ExtraTreesClassifier(bootstrap = True, criterion= 'entropy', max_depth = 6, min_samples_split = 3, n_estimators = 50)  
final_rf.fit(x_train,y_train)  
final_pred = final_rf.predict(x_test)  
final_score = accuracy_score(y_test,final_pred)  
final_score
```

```
0.7829503469406146
```


I can see that we have imported “GridSearchCV” and I have used selected parameters and here I use Cross validation “5”, and I train the model and select the best parameters and also I have predicted the final accuracy score which is 94.3%

ROC AUC Curve

```
from sklearn.metrics import roc_curve, auc

fpr, tpr, thresholds = roc_curve(extr_pred, y_test)
roc_auc = auc(fpr, tpr)

plt.figure(figsize = (10,9))
plt.plot(fpr, tpr, lw=5, color = 'cyan', label = 'ROC Curve(area = %0.2f)%roc_auc)
plt.plot([0,1],[0,1],lw =5, color = 'black', linestyle = '--')
plt.xlim(0.0,1.0)
plt.ylim(0.0,1.0)
plt.xlabel('False positive rate', fontsize = 18)
plt.ylabel('True positive rate', fontsize = 18)
plt.title('ROC AUC Curve', fontsize = 25)
plt.legend(loc = 'lower right', fontsize = 18)
plt.show()
```



```
import joblib
joblib.dump(final, 'Micro_credit.pkl')

['Micro_credit.pkl']
```

Conclusion:

- *I have built a model, I have used multiple models but the highest score that I have received is of Extra Trees Classifier model. So, this is the best model for predicting the values here.*
- *I have made box plot, so from their I come to know that there were a lot of outliers present, So, I have treated them as well as.*
- *In the dataset there was the problem of skewness that I have observed, So I have treated them also.*
- *These are the keys which are used for model prediction of our dataset: -*
 - *Average precision is 0.96, F1 Score is 0.96 and ROC – AUC Score is also 0.958*

Limitations and Scope for the Future:

- *There was Class Imbalance which had to be handled because if we don't do that then our model would become biased, So I have to use resampled functions to treat this thing and there are chances that now it may affect the model.*
- *As there were a lot of outliers and skewness present, so data loss was also there.*

